A risk premium for every policy period, built and challenged with AI

The Per-peril blend · Manuel Rodríguez · individual team · Part 1, established market

0.12996

public score, first among 19 teams on the public leaderboard (19 Oct)

0.138

out-of-sample D^2 on 2024, against 0.071 for a constant rate per line

217

posts by three AI work lines in a shared forum, in 54 hours

Five-page report · offline model dashboard · reproducible code · record of the AI forum and decisions

The task: expected ultimate cost per policy period

626,639 periods in motor, buildings and contents, learned from 2,409,311 periods of 2021–2024. Net of excess and recoveries; declined claims excluded.

Most periods cost nothing

Motor

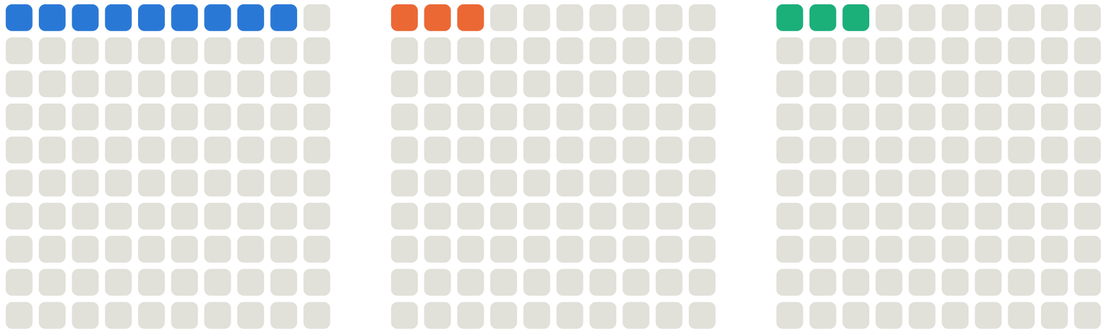
9.1% with a claim cost

Buildings

3.5% with a claim cost

Contents

3.1% with a claim cost

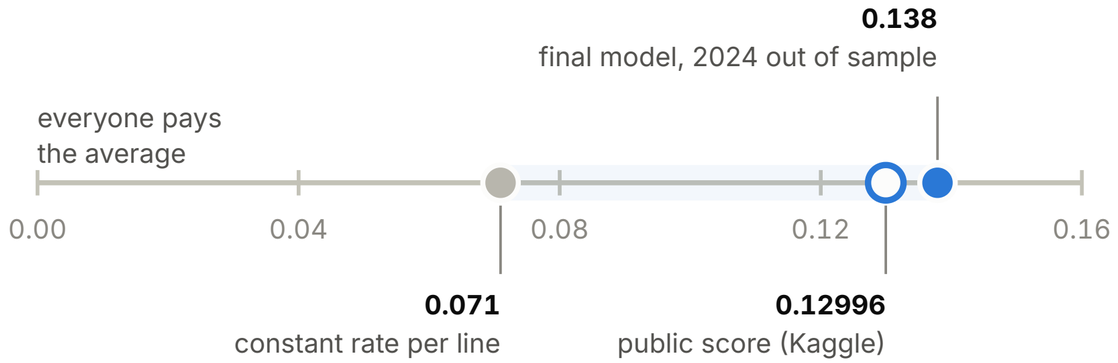


Each square is one policy period in 2024; coloured squares had a claim cost.

How it is scored: Tweedie D^2, power 1.5

D^2 = 1 – deviance of the model / deviance of the average

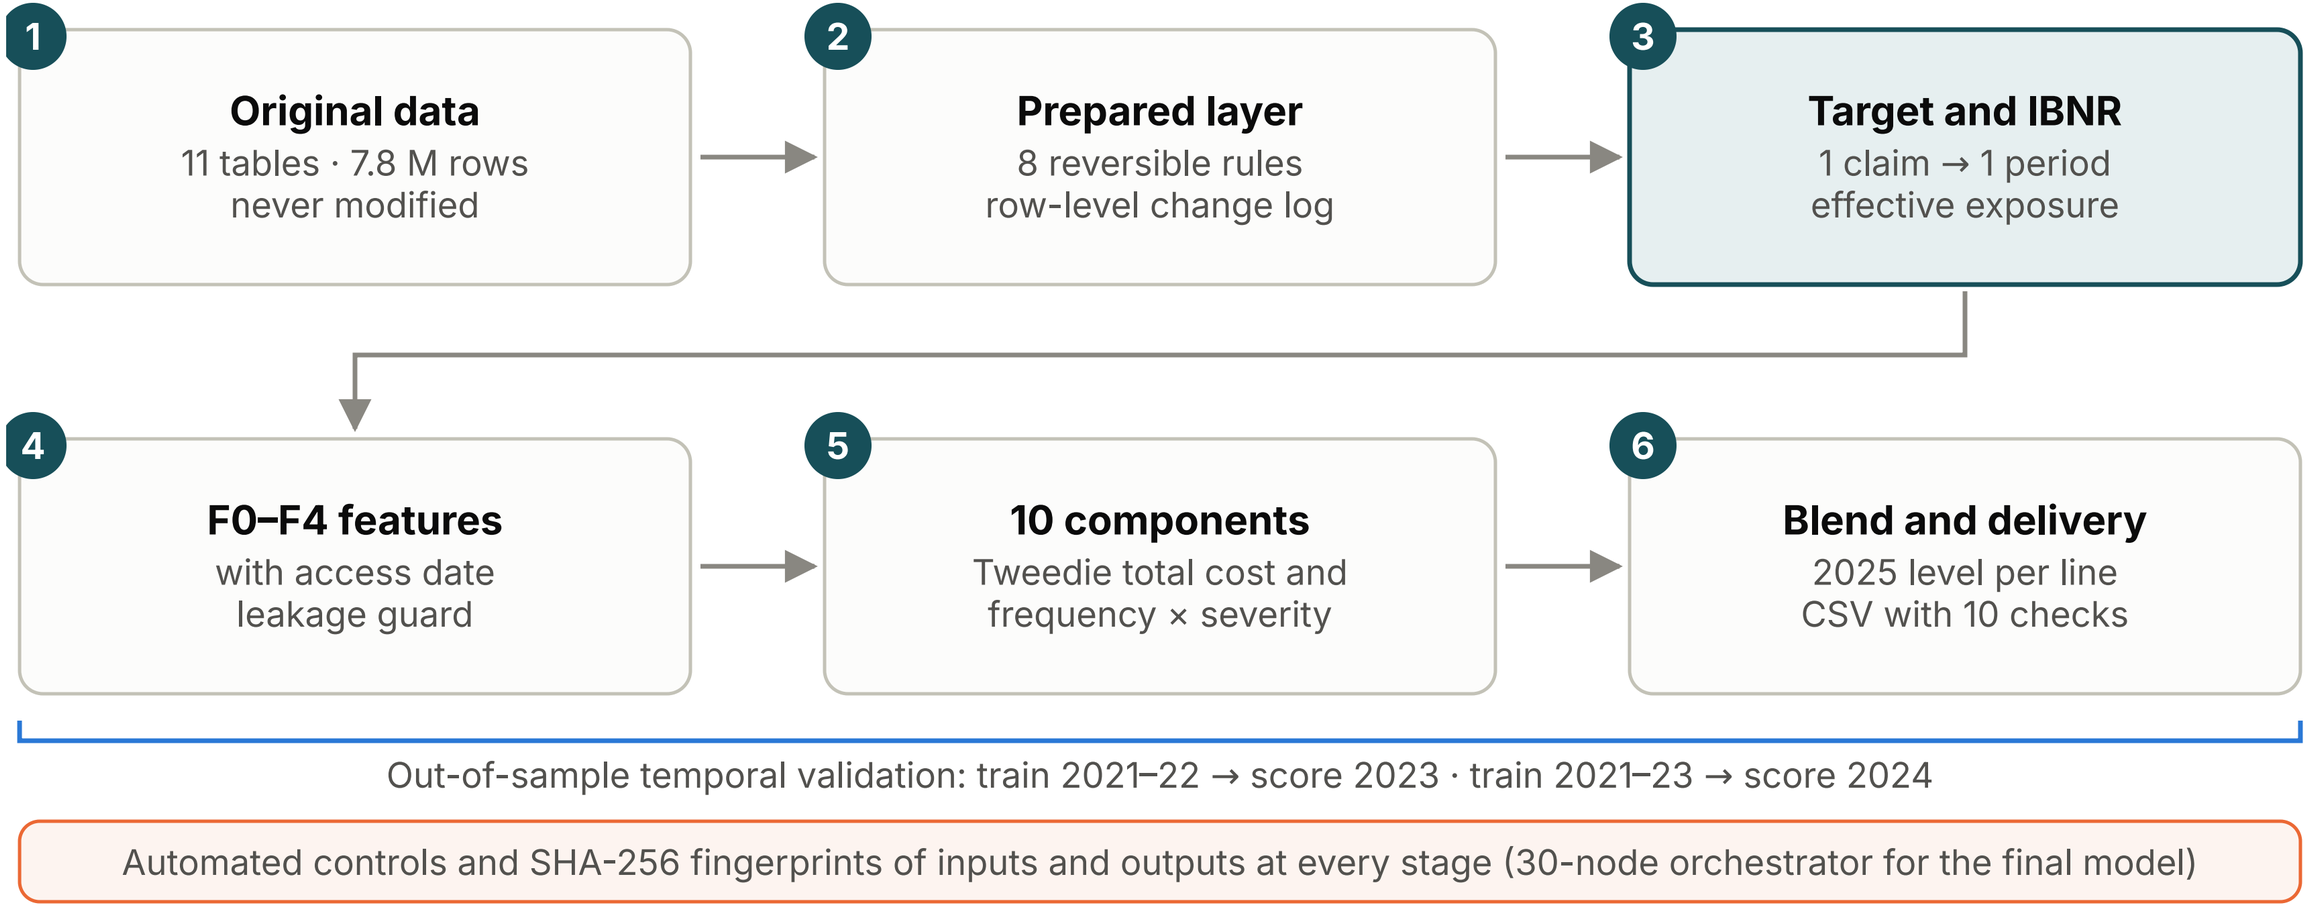
Tweedie deviance with power 1.5, summed over all periods



The scale reaches 1 only if every period's cost were predicted exactly.

Six stages, each one checked

The reporting delay enters once, through the training exposure. Stages 4 to 6 are validated on full out-of-sample years.



Data: fix what can be proven, declare what cannot

The eleven original files stay untouched. Eight reversible rules keep a row-level log, and an independent rebuild matched 148 million cells.

Problem found	Rows	Treatment
Motor customer IDs had lost their leading zeros	123,664	Padded to eight digits; join restored
Exposure stored in days or on a 360-day basis	14,229	Covered days / 365
Sums insured jumping about ×1,000	2,073 + 1,454	Divided by 1,000; 149 lack historical confirmation
Duplicate claims	604	Copy excluded
Declined renewals, claims before the policy history	12,041 · 413	Audit flags only, never predictors

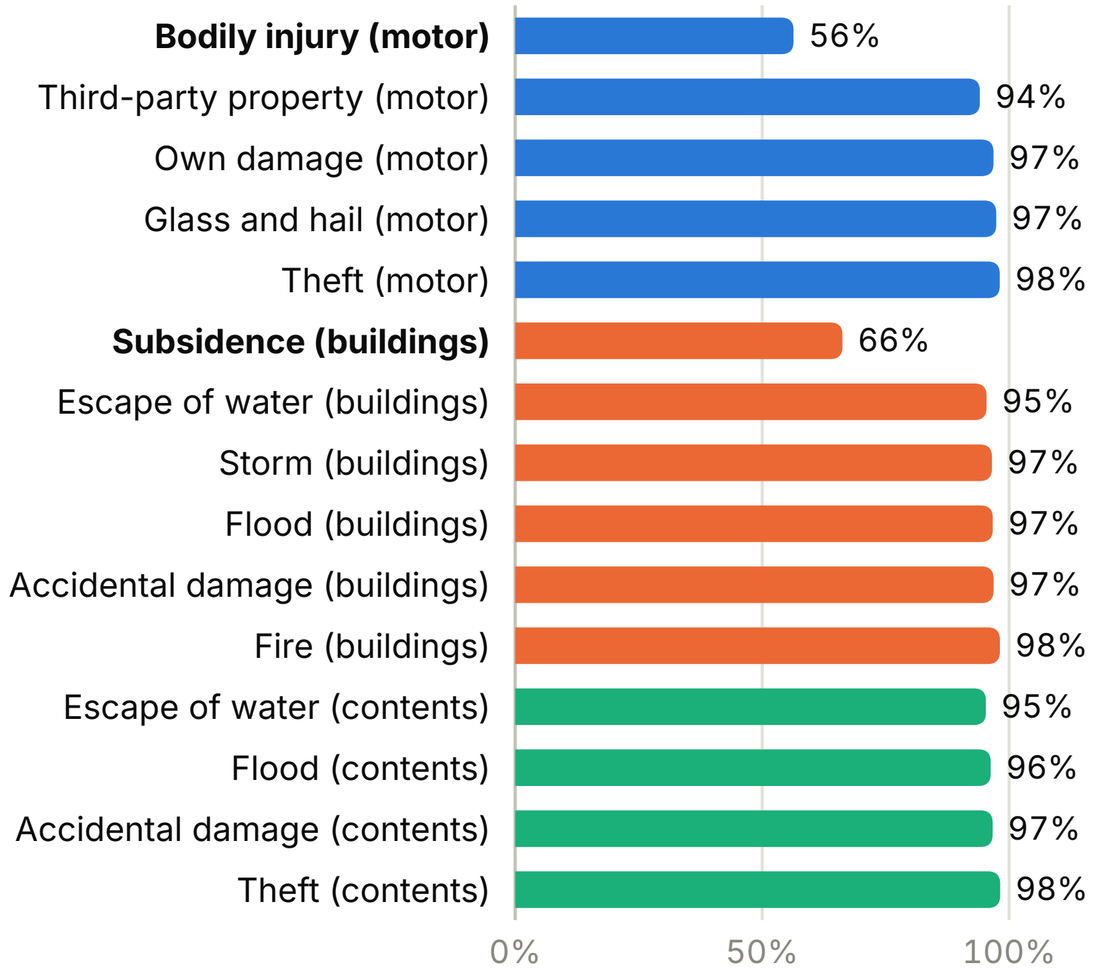


includes 90.2 M CU of open case reserves

Part of 2024's claims was still unreported

A chain-ladder on claim counts per peril estimates how much of the 2024 exposure is already "seen". Two perils stand out.

Share of 2024 exposure already reported, by peril



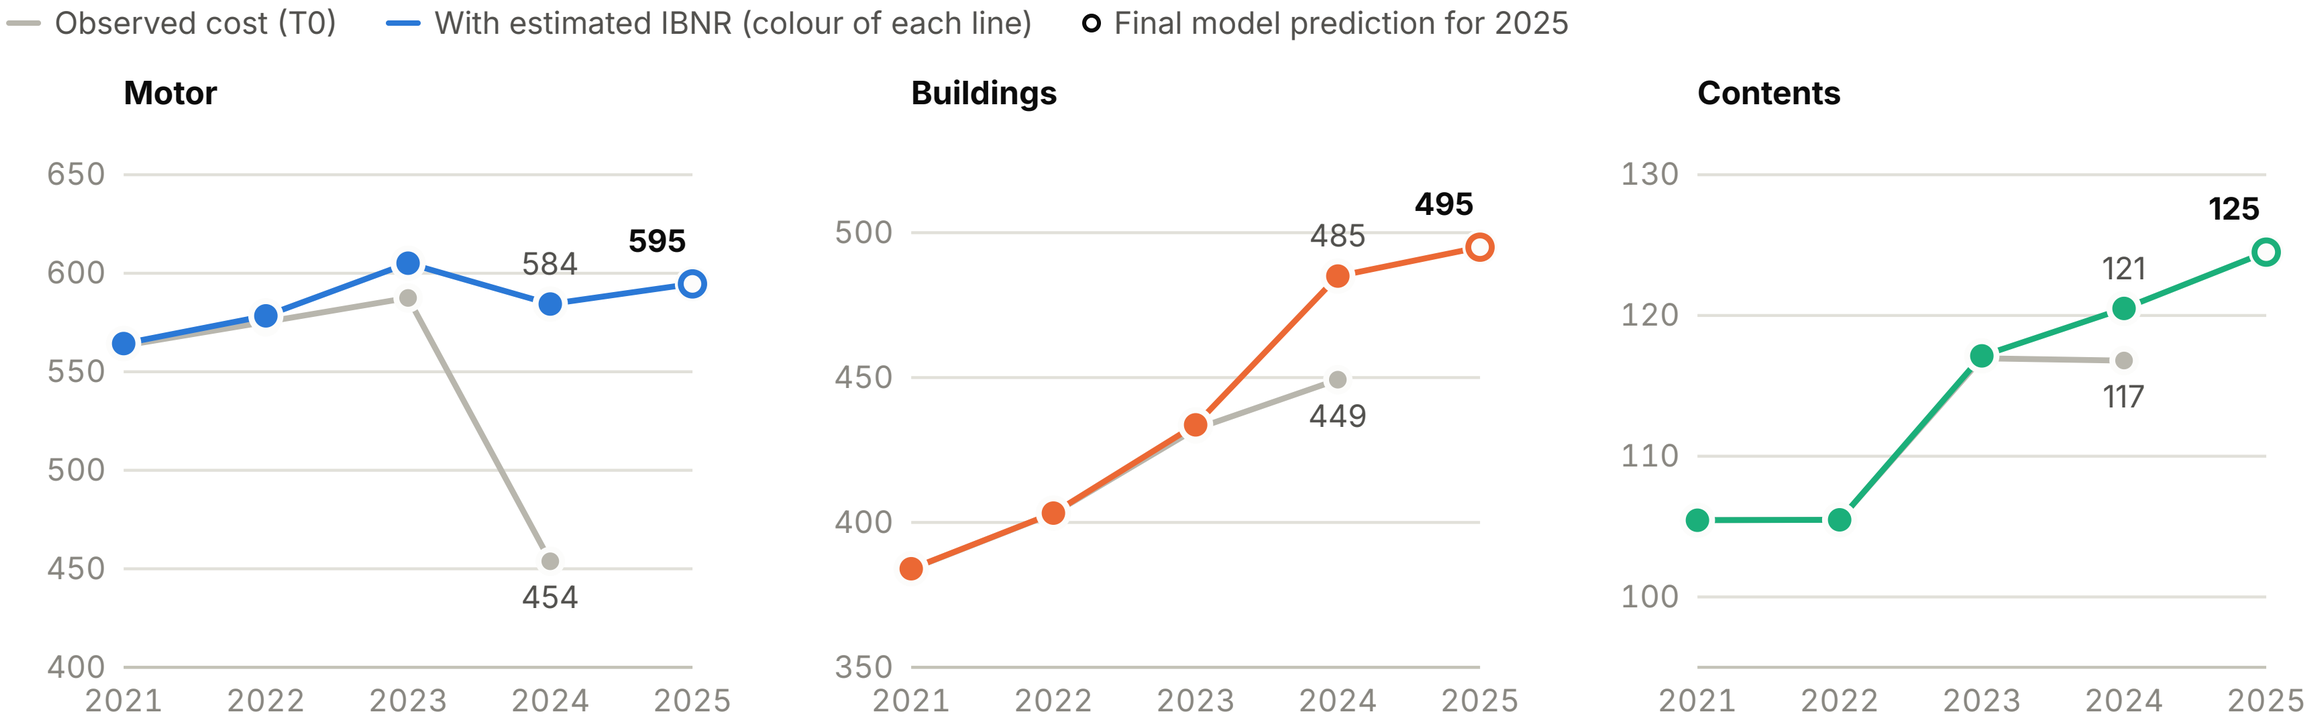
How the delay enters the model

Motor bodily injury, nominal exposure 2024	**1.00**
× share of its claims already reported	**0.56**
= effective exposure, the offset in training	**0.56**
Exposure used to predict 2025	**nominal**

The model learns the full-year rate from the claims seen so far, so IBNR is not added a second time. Building subsidence is at 66%; the other perils at 94–98%.

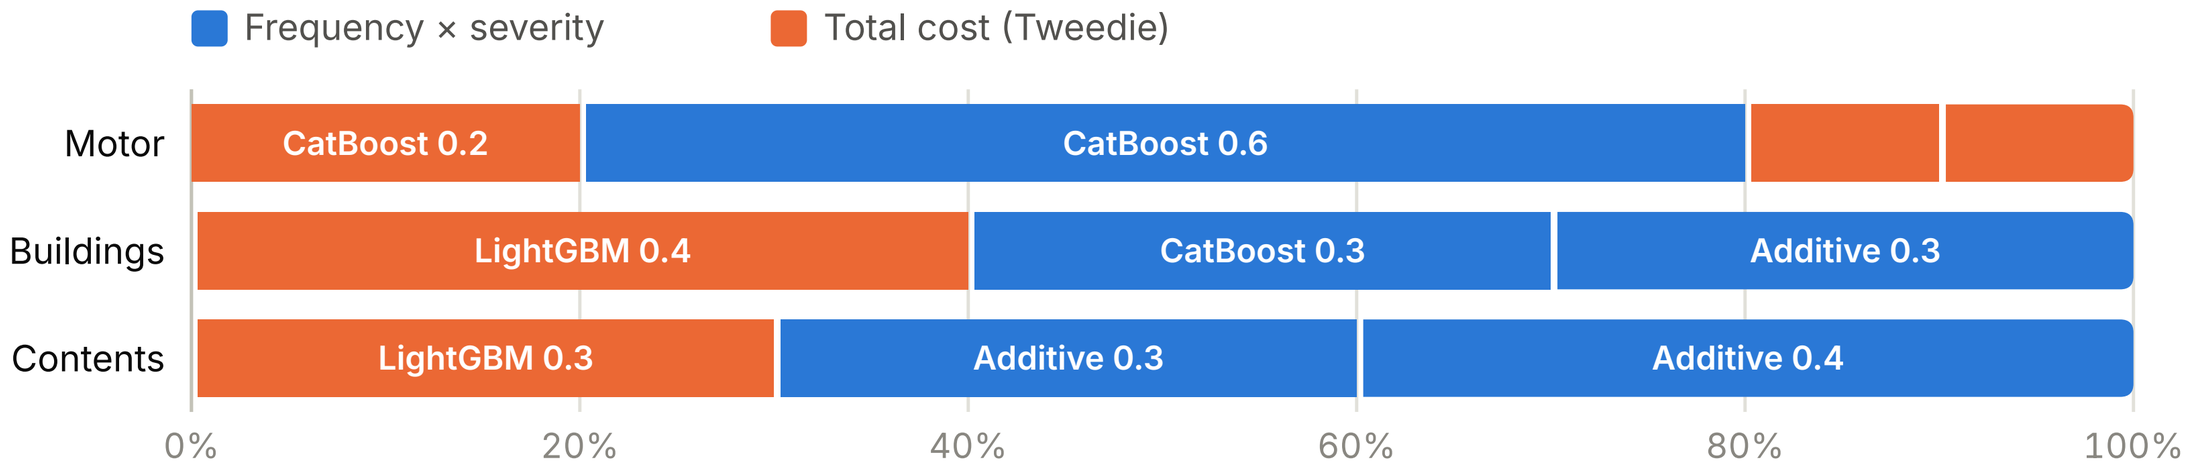
With IBNR, 2024 is no longer an outlier

Motor risk premium per exposure year: 454 CU observed in 2024, 584 with the estimated IBNR, and 595 predicted for 2025.



The Per-peril blend: one model, ten components

Each line blends Tweedie total-cost models with per-peril frequency × severity. Changes were judged on 2023 and 2024, with all losses and without the most extreme ones, under criteria written before computing.



594.6 CU

Motor: mean 2025 risk premium per exposure year

Level factor 0.996 (2023: 0.986 · 2024: 1.005)

Severity trend per peril, capped at ±15%

495.0 CU

Buildings: mean 2025 risk premium per exposure year

Level factor 1.097 (2023: 1.059 · 2024: 1.134)

Severity trend per peril, capped at ±15%

124.5 CU

Contents: mean 2025 risk premium per exposure year

Level factor 1.085 (2023: 1.078 · 2024: 1.092)

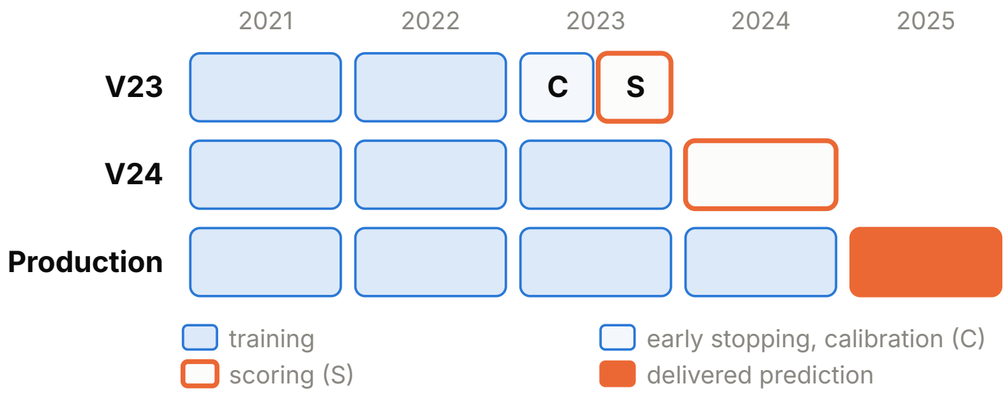
No severity trend

2025 level per line: f = Σ(y/√μ) / Σ√μ, the factor that minimises the Tweedie deviance, averaged over the 2023 and 2024 cuts.

Fitted on past years, scored on a later full year

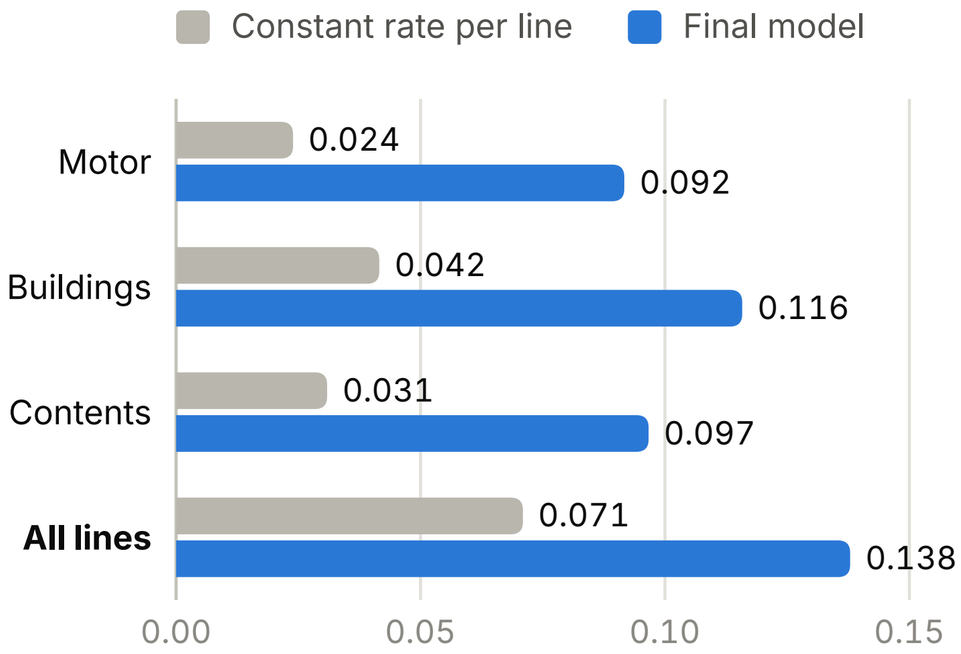
On 2024 the final model explains almost twice the deviance of a constant rate per line (0.138 against 0.071) and improves every line.

Three temporal cuts



C and S: 2023 customer halves used to calibrate and to score.
Labels valued at 31-12-2024: an approximate temporal validation.

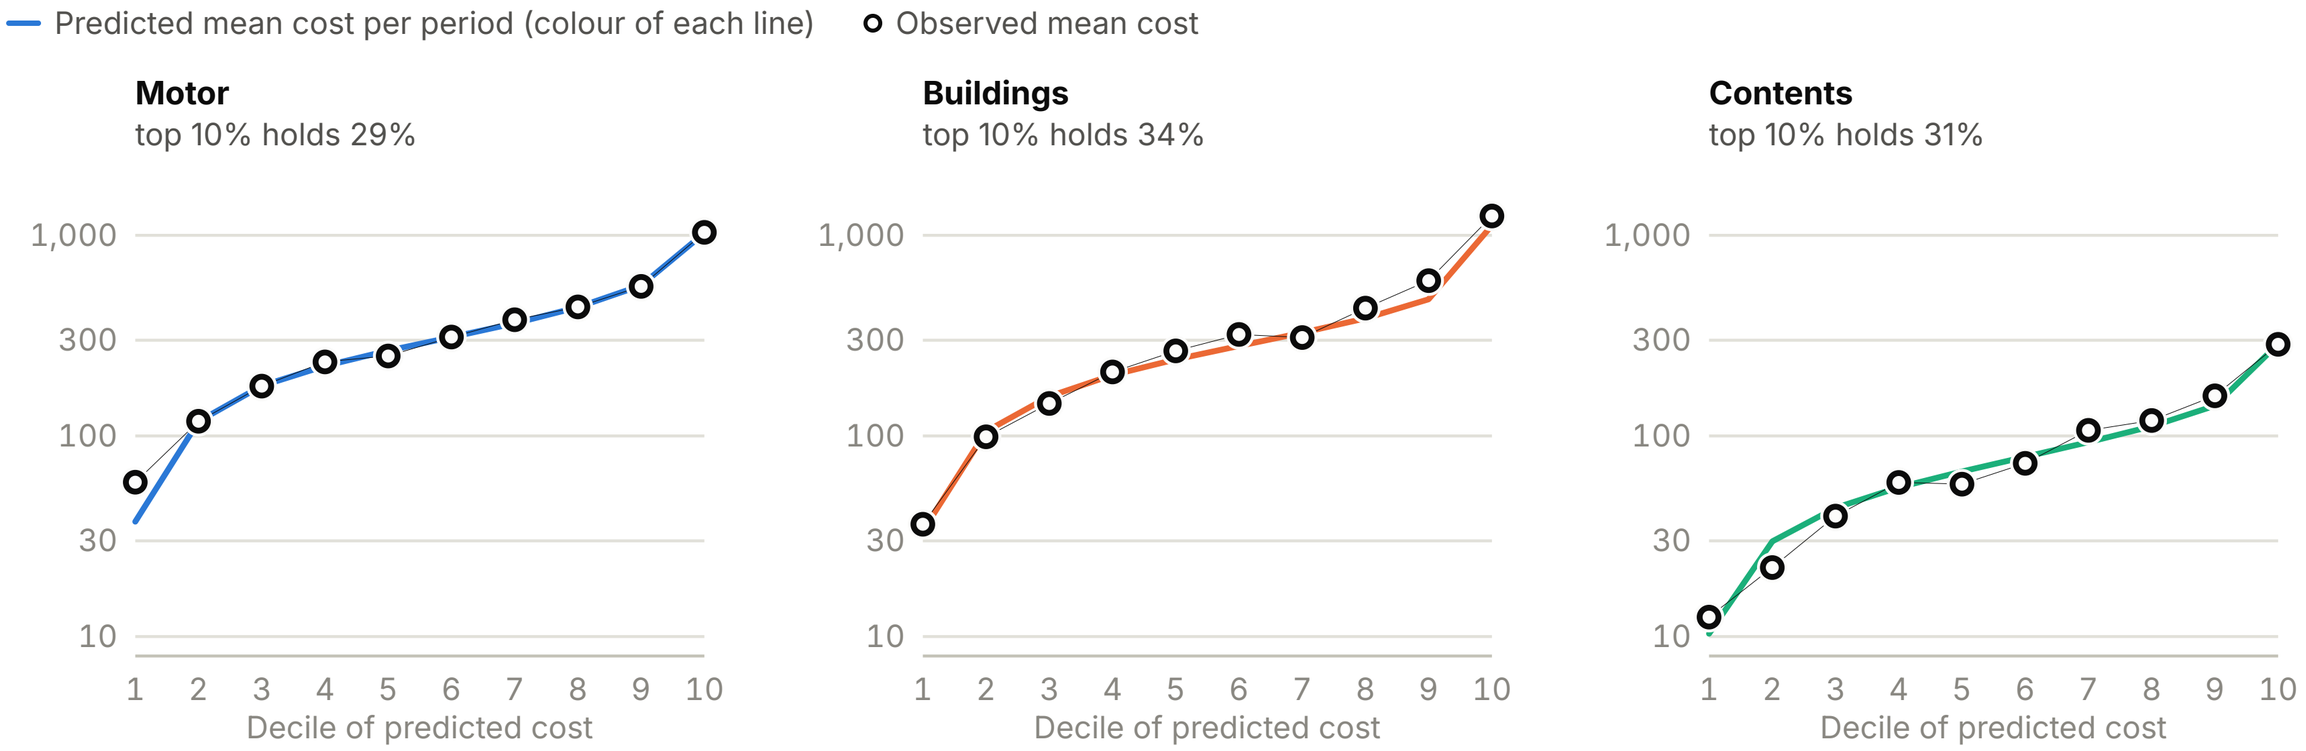
2024 D^2 by line



After the choice, eleven pre-registered experiments, an IBNR backtest among them, tested the model again; the report lists each one with its result.

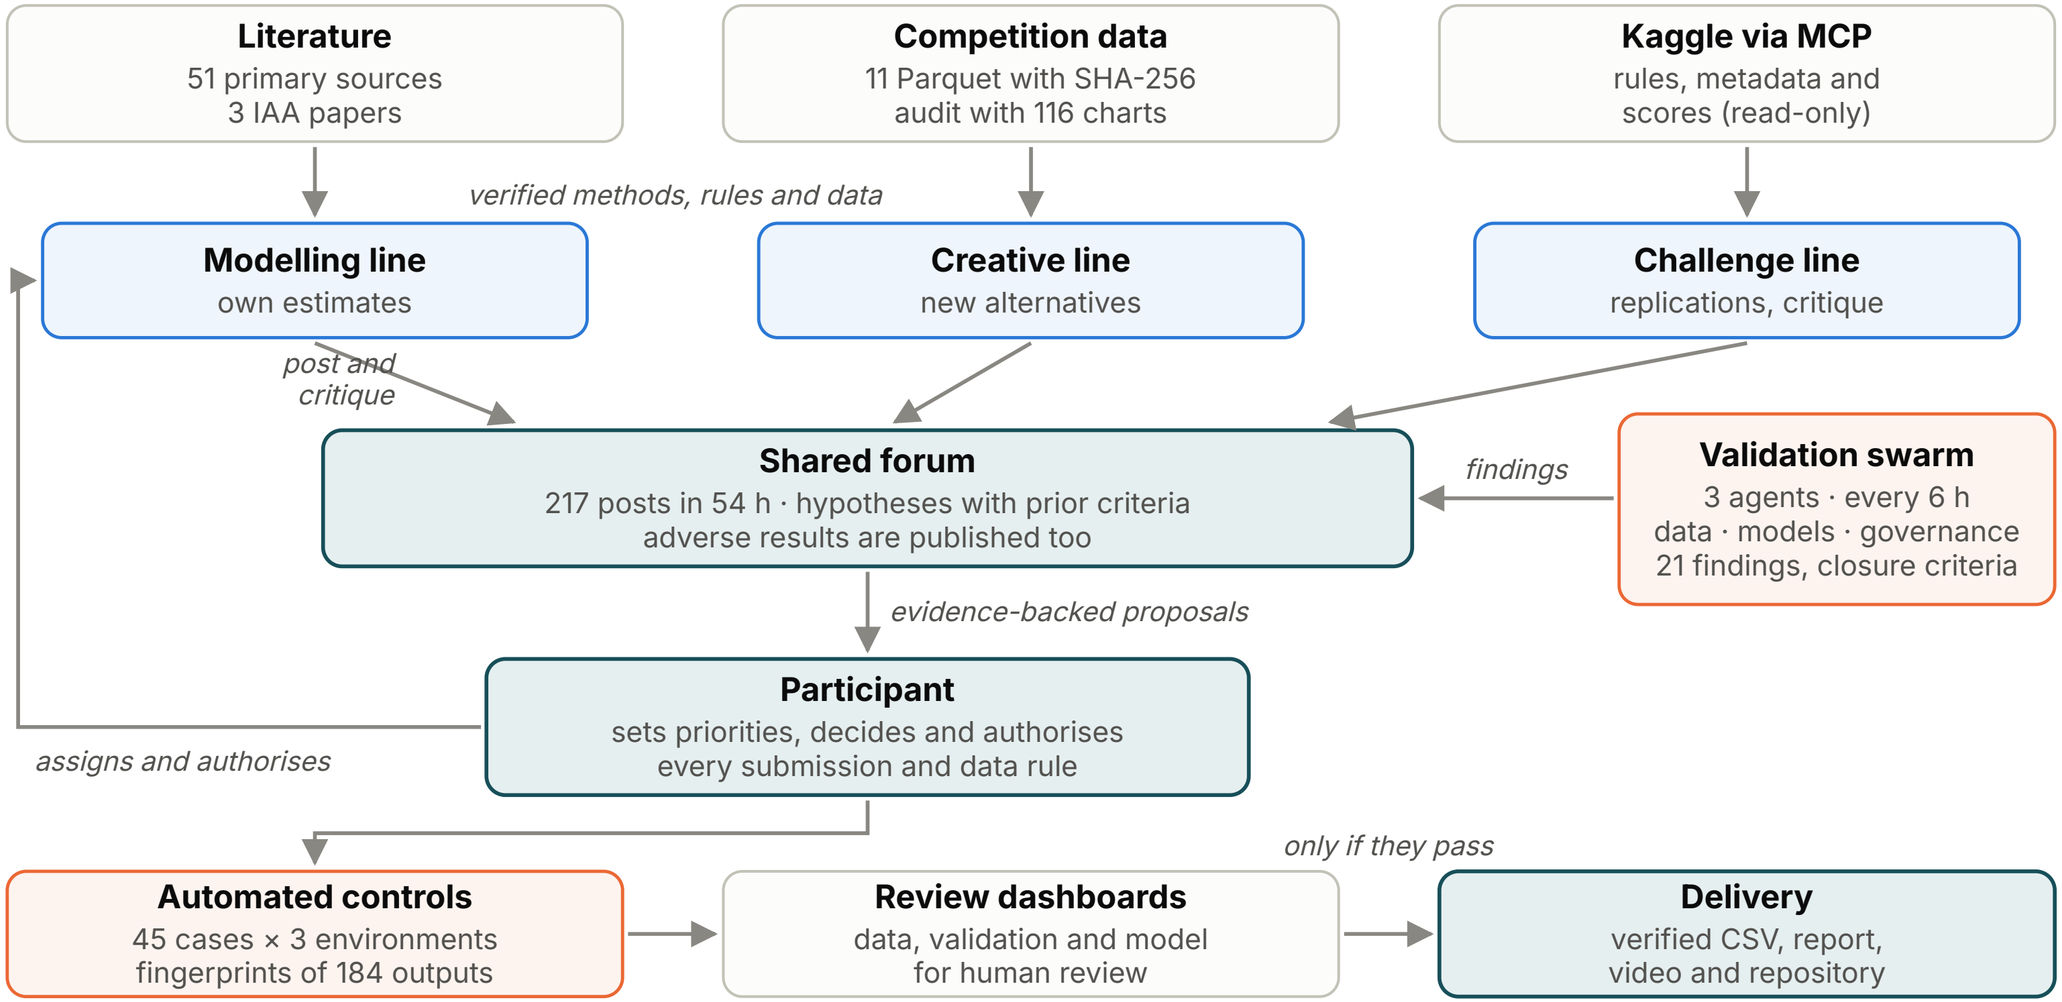
It ranks risk and stays calibrated

The 10% of periods with the highest predicted cost hold 29–34% of the observed 2024 cost. With case reserves ×0.75 or ×1.25, D^2 only moves from 0.1379 to 0.1372, while the 2025 level moves up to 6.8%.



AI was organised so that it never validated itself

Two assistants, Claude and Codex, across three independent lines. Hypotheses were posted with their acceptance criterion before computing, and adverse results were published too.



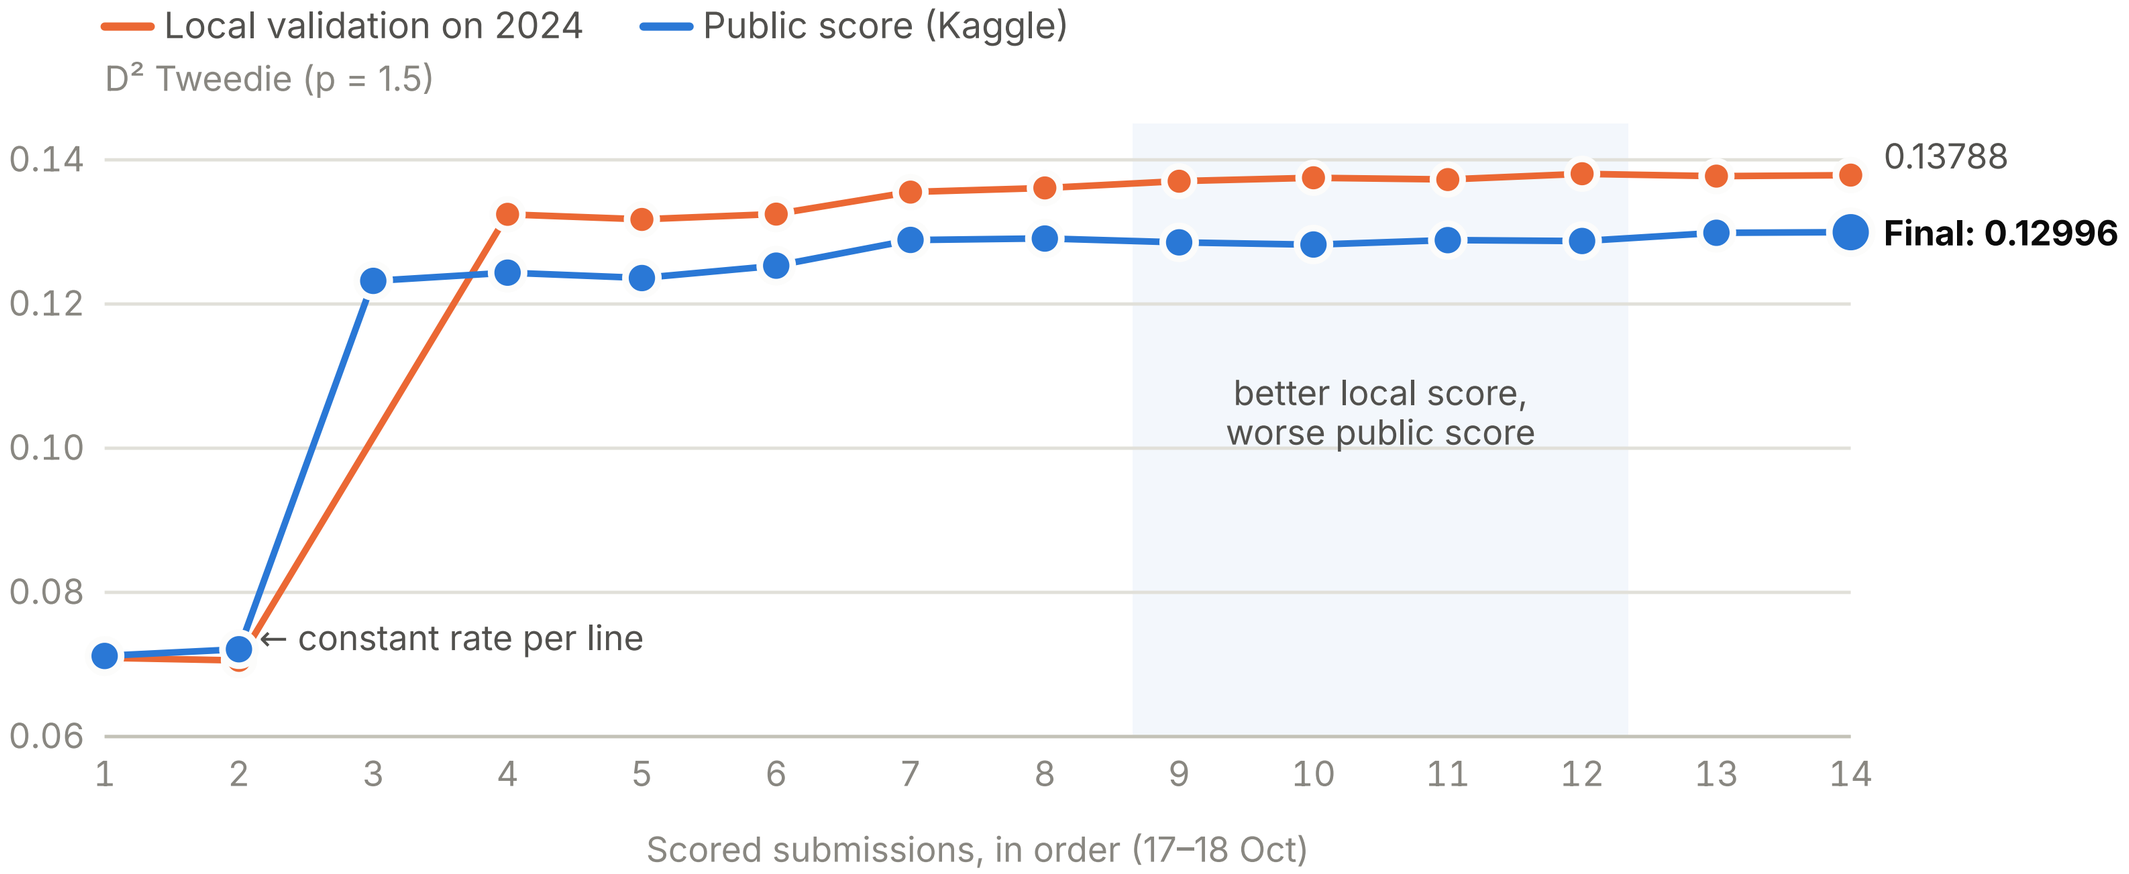
What AI did at each stage, and how it was checked

Every stage had its own check, and the decisions that reached the model went through my authorisation.

Stage	What AI did	How it was verified
Data and target	Audited 105 columns with 174 checks; found the missing zeros, exposure in days, sums ×1,000 and 604 duplicates	Independent recomputation (41/41) and byte-identical rebuild
Features	Availability contract per column; flagged a variable that measures the period's own claim (D^2 0.43; all others below 0.04)	Excluded in code by a predictor guard
Methodology	Two literature searches: 51 primary sources checked against DOI and arXiv, plus three IAA papers	Unsupported claims dropped, such as a 17% gain proposed by another AI tool
Competition context	Rules, metadata and scores read from Kaggle through an MCP connection, read-only	Every submission receipt kept with its CSV fingerprint
Modelling	44 numbered contrasts in the challenge line and pre-registrations in the other two	Intervals clustered by customer, address and region; reserve scenarios
Testing and governance	30-node orchestrator with fingerprints of 184 outputs; validation swarm	A validator found the package could report success with missing files: fixed and tested with 36 adverse cases

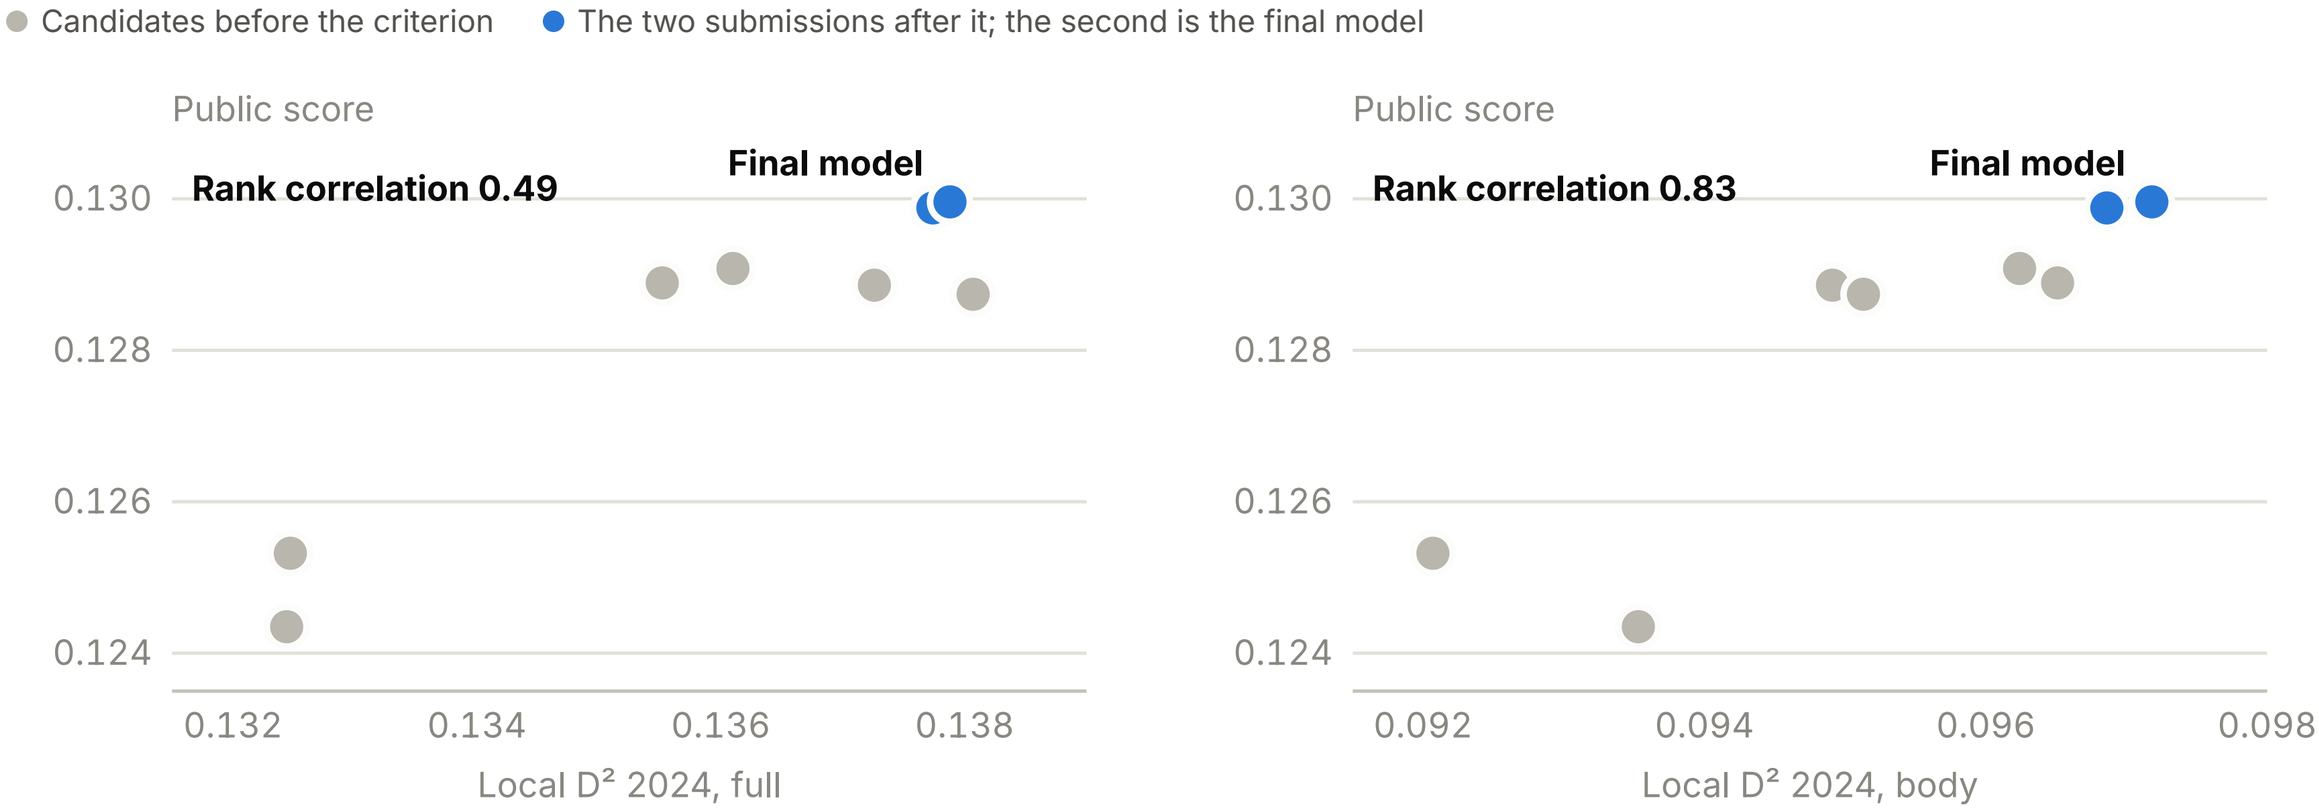
Local gains that the public score did not confirm

The 14 scored submissions in order. In the shaded zone the local 2024 score kept rising while the public score fell; the forum set out to explain why.



The forum's answer: judge the body, keep the tail

A few extreme claims drove the local score. Without the costliest 0.5% of periods, local D² ranks candidates like the public score (rank correlation 0.83 against 0.49). The two submissions that followed set the public maximum.



A swarm of agents audited the project every six hours

Three agents, for data, models and governance, raised 21 findings. Each one states how it can be closed and which human reviewer it needs.

Findings by area and state



High severity · Closed by my decision, technical scope only

Package controls allowed success despite non-conformities

Fixed in code and tested with 36 adverse cases; independence and reserves stay as declared limitations.

High severity · Documents corrected and verified

Overstated causal exclusion of reserve inadequacy

The claim about reserves was downgraded to a hypothesis.

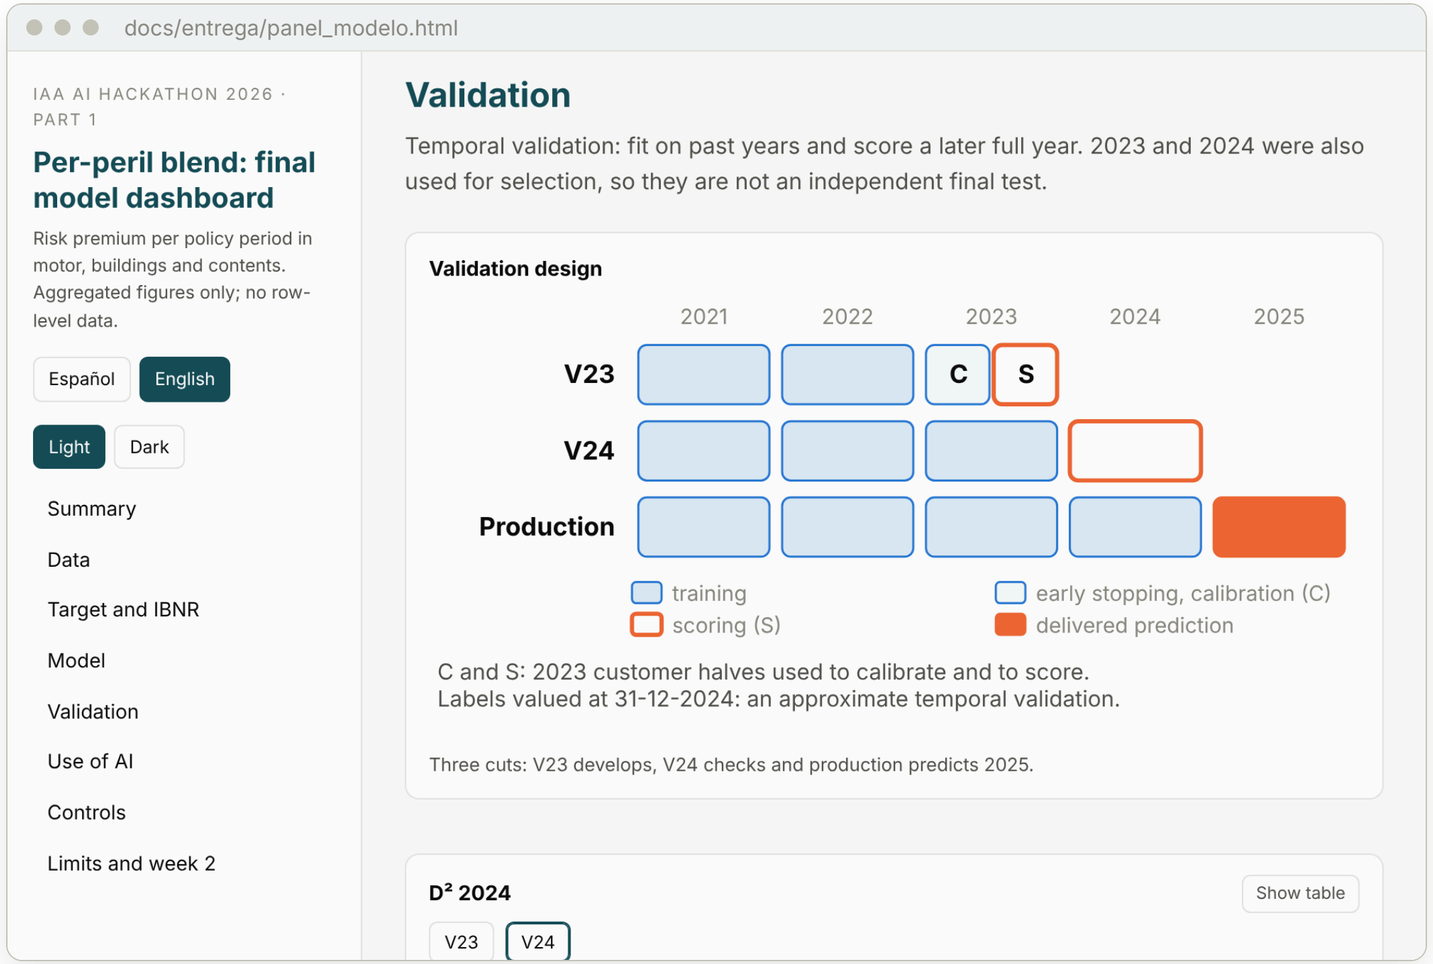
High severity · Open

Availability at quotation not proven for conditional features and annual weather

Needs the data owner to confirm the dates.

A dashboard to check the numbers without reading code

An offline HTML file built from the same data file as the report, which records the path and SHA-256 of its 48 sources.



- Eight sections, from data and target to validation, AI use, controls and limits.
- Every chart switches to the table behind it; the swarm's findings filter by area.
- English or Spanish, in light or dark mode.
- Kept local and not published online, because it holds aggregates of competition data.

The delivered file can be rebuilt and checked

45 of 45 control cases pass in the cloud, on a virtual machine and on Windows. The final CSV is rebuilt from saved components with the SHA-256 that Kaggle received.

Layer	Result
Original data and rebuild	11 of 11 identical; 0 differences in 148 M cells
Target and exposure	18 of 18 checks
Features	25 of 25 checks
Package: 9 nominal and 36 adverse cases	45 of 45 in three environments
Delivered CSV	10 of 10 checks

```
$ python reproducir.py mezcla --objetivo=S30

receipt Windows · 19 Oct 2026
final CSV same SHA-256 as the file Kaggle received
intermediate max relative difference 7.2e-16
exit code 3 = CSV identical, tolerated intermediate differences

$ python reproducir.py todo --objetivo=S30

receipt reference cloud · 20 Oct 2026 · 36 nodes
final CSV same SHA-256 as the file Kaggle received
intermediate 162 of 181 equal; max. 3.9e-15
replay 436 estimators, no refit: 194 of 194 outputs identical
```

What we cannot prove yet, and how Week 2 will test it

Eleven pre-registered experiments, an IBNR backtest among them, measured these limits; the report gives each one its effect or the check that would close it.

Limits we declare

- With one claims valuation, reserve adequacy cannot be proven: a 3.7% error in motor reserves already moves the level 1%.
- The 2025 level is the weak point: ±3.4% in buildings, and one scenario puts motor 1.5% to 2.5% high.
- 2023 and 2024 chose the model, and its lead over the previous best submission is not robust.
- Feature dates at quotation still need the data owner's confirmation.
- 149 divided sums insured weigh 0.01% of the total but decide their own rows.
- Retraining and replay were run by the same workflow, not by an independent operator.

Week 2 protocol for the new market

What we will compare	Planned response
Contract, units and exposure	Rebuild the prepared layer
Peril mix and tails	Refit severities; review the body criterion
Claims development	Re-estimate the chain-ladder
Available features	Update the contract and the predictor guard
Frozen model performance	Recalibrate first; rebuild only if needed

In the repository: docs/entrega/ (report, dashboard, slides) · reproducir.py · docs/expediente/ (forum, decisions and receipts)